

Homework 1

Ian Frankenburg


4/7/2020

Bayesian Adaptive Lasso

Part a.

Consider $p = 1$. Simulate 5,000 Monte Carlo samples from the conditional prior $\beta|\tau^2 = 1$ and obtain a plot of the density using the R function density.


```
n <- 5000  
plot(density(rnorm(n,0,1)), main=TeX(paste("$\\beta$", "marginal")))
```



Part b.

Consider $p = 1$. Simulate 5,000 Monte Carlo samples from the marginal prior β , considering $\lambda^2 = 2$, so that $\mathbb{E}(\tau^2|\lambda) = 1$. Obtain a plot of the density as in **a**.

```
lambda <- sqrt(2)
tau.sq <- rgamma(n,shape=1,rate = lambda^2/2)
beta.marginal <- rnorm(n,0,sqrt(tau.sq))
plot(density(beta.marginal), main=TeX(paste("\\lambda^2 = 2$")), xlim=c(-5,5))
```



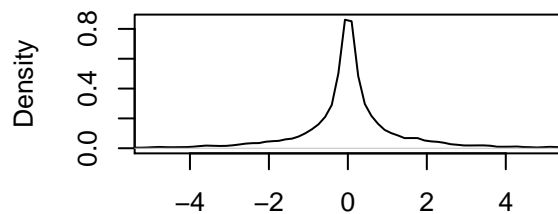
N = 5000 Bandwidth = 0.1185

Part c.

Consider $p = 1$. Add a hyper prior on $\gamma = 1/\gamma \sim \text{Gamma}(a, \text{rate} = b)$. Assess how the marginal prior of β changes for $a = 1$ and values of $b \geq 1$.

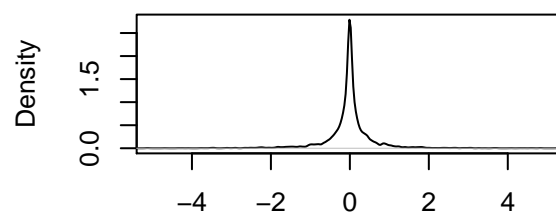
```
set.seed(1)
par(mfrow=c(2,2))
rates <- c(1,3,5,10)
for(b in rates){
  lambda <- 1/rgamma(n,1,b)
  tau.sq <- rgamma(n,shape=1,rate = lambda^2/2)
  beta.marginal <- rnorm(n,0,sqrt(tau.sq))
  plot(density(beta.marginal), main=paste("rate b = ",b),xlim=c(-5,5))
}
```

rate b = 1



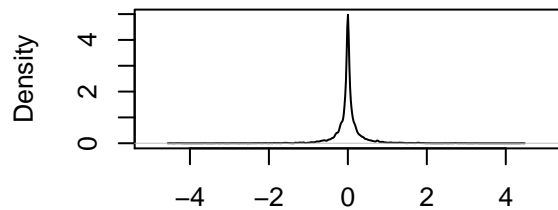
N = 5000 Bandwidth = 0.09665

rate b = 3



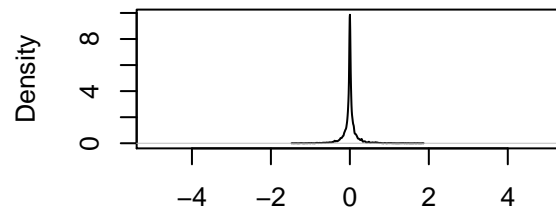
N = 5000 Bandwidth = 0.03314

rate b = 5



N = 5000 Bandwidth = 0.01929

rate b = 10



N = 5000 Bandwidth = 0.009661

Part d.

Considering the hyper prior in **c.**, describe a Markov Chain Monte Carlo algorithm to sample from the posterior distribution of β and σ^2 .

I will implement a joint Gibbs and Metropolis sampler. The model is

$$\begin{aligned} \mathbf{Y}|\beta, \sigma^2 &\sim N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) \\ \beta_j|\tau_j^2 &\sim N(0, \tau_j^2) \\ \tau_j^2 &\sim \text{Gamma}(1, \frac{\lambda^2}{2}) \\ \lambda^2 &\sim \text{Inverse-Gamma}(a, 1/b) \\ \sigma^2 &\sim \text{Inverse-Gamma}(0.1, 0.1). \end{aligned}$$

I need the full conditionals

$$\begin{aligned} &\{\beta_1, \dots, \beta_p | \mathbf{Y}, \sigma^2, \tau_1^2, \dots, \tau_p^2, \lambda\}, \\ &\{\sigma^2 | \mathbf{Y}, \beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \lambda\}, \\ &\{\tau_1^2, \dots, \tau_p^2 | \mathbf{Y}, \beta_1, \dots, \beta_p, \sigma^2, \lambda\}, \\ &\{\lambda | \mathbf{Y}, \beta_1, \dots, \beta_p, \sigma^2, \tau_1^2, \dots, \tau_p^2\} \end{aligned}$$

which are all proportional to

$$p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) \times p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) \times p(\tau_1^2, \dots, \tau_p^2 | \lambda) p(\lambda) p(\sigma^2)$$

so I'll start with the posterior

$$\begin{aligned} p(\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda | \mathbf{Y}) &\propto p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) \\ &\quad \times p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) \\ &\quad \times p(\tau_1^2, \dots, \tau_p^2 | \lambda) p(\lambda) p(\sigma^2). \end{aligned}$$

As a function of just σ^2 , this is proportional to

$$\begin{aligned} &p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) p(\sigma^2) \\ &= N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) IG(a, b). \end{aligned}$$

Time to show this is inverse-gamma distributed.

$$\begin{aligned} &N(\mathbf{y}; \mathbf{X}\beta, \sigma^2 \mathbf{I}) IG(\sigma^2; q, r) \\ &\propto (\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) \right\} (\sigma^2)^{q-1} \exp \left\{ -\frac{r}{\sigma^2} \right\} \\ &= (\sigma^2)^{-(n/2+q)-1} \exp \left\{ -\frac{1}{\sigma^2} \left(r + \frac{1}{2} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) \right) \right\} \\ &= IG(n/2 + q, r + (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)/2) |_{q=0.1, r=0.1} \end{aligned}$$

As a function of β , the conditional is proportional to

$$\begin{aligned} &p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) \\ &= N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) \cdot \prod_{i=1}^p N(0, \tau_i^2) \\ &= N(\mathbf{X}\beta, \Sigma) \cdot N(0, \Omega), \text{ where } \Omega = \text{diag}(\tau_1^2, \dots, \tau_p^2) \\ &= N(m, \mathbf{M}) \end{aligned}$$

with $m = \mathbf{M}\mathbf{X}^\top \Sigma^{-1} \mathbf{y}$ and $\mathbf{M} = (\mathbf{X}^\top \Sigma^{-1} \mathbf{X} + \Omega^{-1})^{-1}$. As a function of $\tau_1^2, \dots, \tau_p^2$, the conditional is non-standard, but it's proportional to

$$\begin{aligned} & p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) p(\tau_1^2, \dots, \tau_p^2 | \lambda) \\ &= \prod_{i=1}^p N(0, \tau_i^2) \cdot \prod_{i=1}^p IG(1, \frac{\lambda^2}{2}) \end{aligned}$$

Finally, as a function of λ , it's proportional to

$$\begin{aligned} & p(\tau_1^2, \dots, \tau_p^2 | \lambda) p(\lambda) \\ &= \prod_{i=1}^p IG(1, \frac{\lambda^2}{2}) \cdot IG(0.1, 0.1) \end{aligned}$$

Now I can build an algorithm to iteratively update through these conditional distributions. I take the starting value of $\beta^{(0)}$ to be the least-squares solution $\hat{\beta}$.

Result: Samples from joint posterior $p(\beta, \sigma^2 | \mathbf{y})$

for s in # samples **do**

```

     $\sigma^{2(s+1)} \sim IG(n/2 + a, 2b + (\mathbf{y} - \mathbf{X}\beta^{(s)})^\top (\mathbf{y} - \mathbf{X}\beta^{(s)})/2)$ 
     $\beta^s \sim N_p(m, \mathbf{M})$ 
     $m^\top = \mathbf{M} = \tau_1^2, \dots, \tau_p^2 \sim N(\boldsymbol{\tau}^{(s)}, \delta \mathbf{I})^2$ 
     $\lambda^* \sim N(\lambda.s, \delta)^2$ 
     $\log(r) = \log p(\mathbf{y} | \beta^*, \tau_1^2, \dots, \tau_p^2, \sigma^{2(s)}) + \log p(\beta^* | \tau_1^2, \dots, \tau_p^2)$ 
     $\quad - \log p(\mathbf{y} | \beta^{(s)}, \tau_1^2, \dots, \tau_p^2, \sigma^{2(s)}) - \log p(\beta^{(s)} | \tau_1^2, \dots, \tau_p^2)$ 
     $u \sim Unif(0, 1)$ 
    if  $\log(u) < \log(r)$  then
    |    $\boldsymbol{\tau}^{(s+1)} = \beta^*$ 
    else
    |    $\boldsymbol{\tau}^{(s+1)} = \beta^{(s)}$ 
    end

```

end

Algorithm 1: Gibbs and Metropolis

Part e.

Implement such algorithm in R and compare your results with estimates obtained using `glmnet()`. In particular, you should test your results on the diabetes data available from lars, (use the matrix of predictors `x`).

```
loglambda.target <- function(lambda, a, b, tau2){
  (-a - 1)*log(lambda) - (lambda^2/2 * sum(tau2)) - 1/(b*lambda)
}
logtau2j.target <- function(tau2j, lambda, betaj){
  -log(sqrt(tau2j)) - 1/2*(1/tau2j*betaj^2 + lambda^2*tau2j)
}

set.seed(1)
data("diabetes")
X <- cbind(diabetes$x); y <- diabetes$y; n <- nrow(X)
X <- cbind(rep(1,n),X)
p <- ncol(X)
samples <- 50000;
#lambda <- 1e-10
lambda <- 1
tau2 <- rep(1000,p)
beta <- solve(t(X)%*%X)%*%t(X)%*%y
sigma2.keep <- lambda.keep <- rep(0,samples)
beta.keep <- tau2.keep <- matrix(NA,nrow=p,ncol=samples)
sigma2 <- var(y)
sigma2.keep[1] <- sigma2
beta.keep[,1] <- solve(t(X)%*%X)%*%t(X)%*%y
lambda.keep[1] <- lambda
tau2.keep[,1] <- tau2
s <- 2
post.means <- matrix(NA,nrow=p,ncol=length(lambdas))
index <- 1
b<-1
for(s in 2:samples){
  lambda.p <- exp(log(lambda)+rnorm(1,0,lambda))
  logr <- loglambda.target(lambda.p,1,b,tau2)-loglambda.target(lambda,1,b,tau2)+log(lambda.p)-log(lambda)
  if(log(runif(1))<logr){lambda <- lambda.p}
  lambda.keep[s] <- lambda
  for(j in 1:length(tau2)){
    tau2j.p <- exp(log(tau2[j])+rnorm(1,0,5))
    logr <-
      logtau2j.target(tau2j.p, lambda, beta[j]) -
      logtau2j.target(tau2[j], lambda, beta[j]) +
      log(tau2j.p)-log(tau2[j])
    if(log(runif(1))<logr){tau2[j] <- tau2j.p}
  }
  tau2 <- 1/rinvgauss(n = length(beta), mean = lambda/abs(beta),
    shape = lambda^2)
  sigma2 <- rinvgamma(1,shape=n/2+0.1,rate=(0.1+t(y-X%*%beta)%*%(y-X%*%beta)/2))
  M <- solve(t(X)%*%X*1/sigma2+diag(1/tau2))
  m <- M%*%t(X)%*%y/sigma2
  beta <- t(rmvnorm(n=1,mean=m,sigma=M))
}
```

```

sigma2.keep[s] <- sigma2
beta.keep[,s] <- beta
tau2.keep[,s] <- tau2
}
post.means[,index] <- matrix(rowMeans(beta.keep[,floor(samples/4):samples]),nrow = p)
index <- index +1
}
fit <- glmnet(X, y)
glmnetCoef <- coef(fit,s=1)
plot(fit, xvar="lambda")

fit <- glmnet(X, y)
glmnetCoef <- coef(fit,s=1)
# df <- cbind(beta.ls,colMeans(beta.keep[floor(samples/4):samples,]),glmnetCoef[-1,])
# colnames(df) <- c("Least Square", "Bayes", "Glmnet")
# df
plot(log(lambda.keep[1:samples]), type="l")
beta.ls <- solve(t(X)%*%X)%*%t(X)%*%y
cbind(beta.ls,betaChain)
beta.keep <- t(beta.keep)
tau2.keep <- t(tau2.keep)
plot(sigma2.keep[1:samples],type="l")
plot(log(tau2.keep[,1]),type="l")
plot((tau2.keep[,2]),type="l")
plot((tau2.keep[,3]),type="l")
plot((tau2.keep[,4]),type="l")
plot((tau2.keep[,5]),type="l")
plot((tau2.keep[,6]),type="l")
plot((tau2.keep[,7]),type="l")

```

Part g.

Free λ and carry out a sensitivity analysis assessing the behavior of the posterior distribution of β and σ^2 , as hyper parameters a and b are changed. Explain clearly the rationale you use to assess sensitivity and provide recommendations for the analysis of the diabetes data.

```
loglambda.target <- function(lambda, a, b, tau2){
  (-a - 1)*log(lambda) - (lambda^2/2 * sum(tau2)) - 1/(b*lambda)
}
logtau2j.target <- function(tau2j, lambda, betaj){
  -log(sqrt(tau2j)) - 1/2*(1/tau2j*betaj^2 + lambda^2*tau2j)
}

set.seed(1)
data("diabetes")
X <- cbind(diabetes$x); y <- diabetes$y; n <- nrow(X)
X <- cbind(rep(1,n),X)
p <- ncol(X)
samples <- 50000;
#lambda <- 1e-10
lambda <- 1
tau2 <- rep(1000,p)
beta <- solve(t(X)%*%X)%*%t(X)%*%y
sigma2.keep <- lambda.keep <- rep(0,samples)
beta.keep <- tau2.keep <- matrix(NA,nrow=p,ncol=samples)
sigma2 <- var(y)
sigma2.keep[1] <- sigma2
beta.keep[,1] <- solve(t(X)%*%X)%*%t(X)%*%y
lambda.keep[1] <- lambda
tau2.keep[,1] <- tau2
s <- 2
lambdas <- seq(from=-12, to = -2, length.out = 12)
post.means <- matrix(NA,nrow=p,ncol=length(lambdas))
index <- 1
b<-1
for(l in lambdas){
  for(s in 2:samples){
    lambda.p <- exp(log(lambda)+rnorm(1,0,lambda))
    logr <- loglambda.target(lambda.p,1,b,tau2)-loglambda.target(lambda,1,b,tau2)+log(lambda.p)-log(lambda)
    if(log(runif(1))<logr){lambda <- lambda.p}
    lambda.keep[s] <- lambda
    for(j in 1:length(tau2)){
      tau2j.p <- exp(log(tau2[j])+rnorm(1,0,5))
      logr <-
        logtau2j.target(tau2j.p, lambda, beta[j]) -
        logtau2j.target(tau2[j], lambda, beta[j]) +
        log(tau2j.p)-log(tau2[j])
      if(log(runif(1))<logr){tau2[j] <- tau2j.p}
    }
    tau2 <- 1/rinvgauss(n = length(beta), mean = lambda/abs(beta),
                      shape = lambda^2)
    sigma2 <- rinvgamma(1,shape=n/2+0.1,rate=(0.1+t(y-X)%*%beta)%*%(y-X)%*%beta)/2))
    M <- solve(t(X)%*%X*1/sigma2+diag(1/tau2))
```



```

m <- M%*%t(X)%*%y/sigma2
beta <- t(rmvnorm(n=1,mean=m,sigma=M))
sigma2.keep[s] <- sigma2
beta.keep[,s] <- beta
tau2.keep[,s] <- tau2
}
post.means[,index] <- matrix(rowMeans(beta.keep[,floor(samples/4):samples]),nrow = p)
index <- index +1
}
fit <- glmnet(X, y)
glmnetCoef <- coef(fit,s=1)
plot(fit, xvar="lambda")
betaChain <- colMeans(beta.keep[,1:samples])
plot(lambdas,post.means[1,],ylab="Coefficients",xlab="log lambdas",ylim=c(min(post.means),max(post.means)))
lines(lambdas,post.means[1,])
for(j in 2:10){
points(lambdas,post.means[j,],ylab="Coefficients",xlab="log lambdas", col=j,type="l")
lines(log(lambdas),post.means[j,], col=j)
}

fit <- glmnet(X, y)
glmnetCoef <- coef(fit,s=1)
# df <- cbind(beta.ls,colMeans(beta.keep[floor(samples/4):samples,]),glmnetCoef[-1,])
# colnames(df) <- c("Least Square", "Bayes", "Glmnet")
# df
plot(log(lambda.keep[1:samples]), type="l")
beta.ls <- solve(t(X)%*%X)%*%t(X)%*%y
cbind(beta.ls,betaChain)
beta.keep <- t(beta.keep)
tau2.keep <- t(tau2.keep)
plot(sigma2.keep[1:samples],type="l")
plot(log(tau2.keep[,1]),type="l")
plot((tau2.keep[,2]),type="l")
plot((tau2.keep[,3]),type="l")
plot((tau2.keep[,4]),type="l")
plot((tau2.keep[,5]),type="l")
plot((tau2.keep[,6]),type="l")
plot((tau2.keep[,7]),type="l")

```

Part h.

Implementation and benchmarking in Julia