

Homework 1

Ian Frankenburg

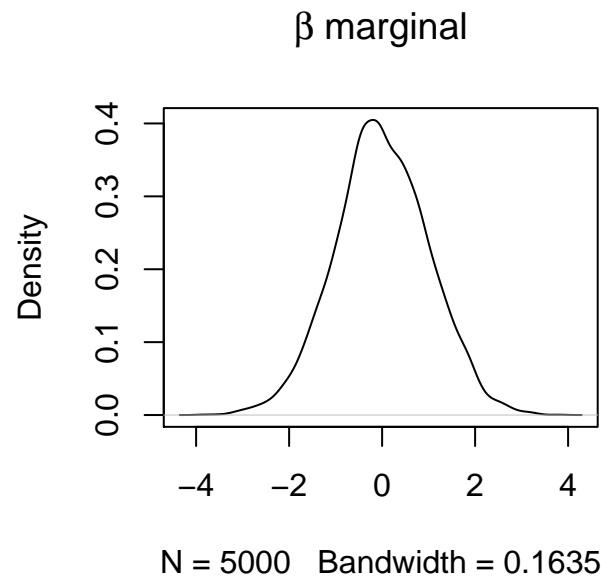
4/7/2020

Bayesian Adaptive Lasso

Part a.

Consider $p = 1$. Simulate 5,000 Monte Carlo samples from the conditional prior $\beta|\tau^2 = 1$ and obtain a plot of the density using the R function density.


```
n <- 5000  
plot(density(rnorm(n,0,1)), main=TeX(paste("$\\beta$", "marginal")))
```



Part b.

Consider $p = 1$. Simulate 5,000 Monte Carlo samples from the marginal prior β , considering $\lambda^2 = 2$, so that $\mathbb{E}(\tau^2|\lambda) = 1$. Obtain a plot of the density as in **a**.

```
lambda <- sqrt(2)
tau.sq <- rgamma(n,shape=1,rate = lambda^2/2)
beta.marginal <- rnorm(n,0,sqrt(tau.sq))
plot(density(beta.marginal), main=TeX(paste("\\lambda^2 = 2$")), xlim=c(-5,5))
```



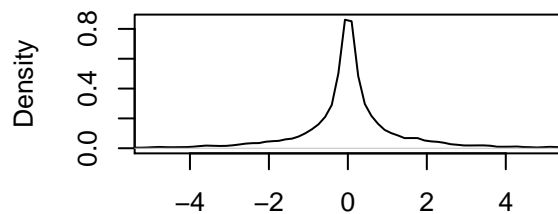
N = 5000 Bandwidth = 0.1158

Part c.

Consider $p = 1$. Add a hyper prior on $\gamma = 1/\gamma \sim \text{Gamma}(a, \text{rate} = b)$. Assess how the marginal prior of β changes for $a = 1$ and values of $b \geq 1$.

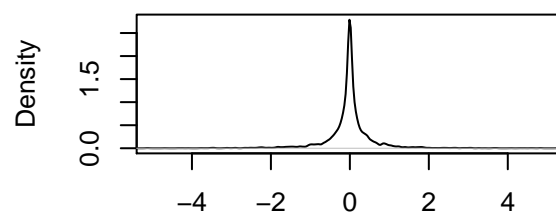
```
set.seed(1)
par(mfrow=c(2,2))
rates <- c(1,3,5,10)
for(b in rates){
  lambda <- 1/rgamma(n,1,b)
  tau.sq <- rgamma(n,shape=1,rate = lambda^2/2)
  beta.marginal <- rnorm(n,0,sqrt(tau.sq))
  plot(density(beta.marginal), main=paste("rate b = ",b),xlim=c(-5,5))
}
```

rate b = 1



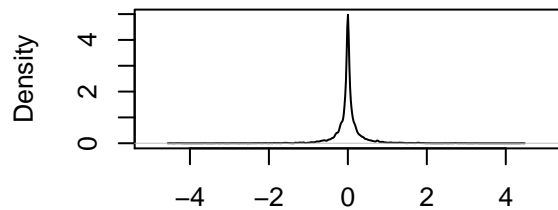
N = 5000 Bandwidth = 0.09665

rate b = 3



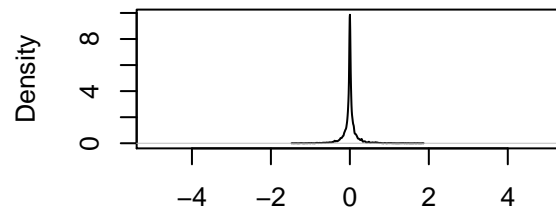
N = 5000 Bandwidth = 0.03314

rate b = 5



N = 5000 Bandwidth = 0.01929

rate b = 10



N = 5000 Bandwidth = 0.009661

Part d.

Considering the hyper prior in **c.**, describe a Markov Chain Monte Carlo algorithm to sample from the posterior distribution of β and σ^2 .

I will implement a joint Gibbs and Metropolis sampler. The model is

$$\begin{aligned} \mathbf{Y}|\beta, \sigma^2 &\sim N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) \\ \beta_j|\tau_j^2 &\sim N(0, \tau_j^2) \\ \tau_j^2 &\sim \text{Gamma}(1, \frac{\lambda^2}{2}) \\ \lambda^2 &\sim \text{Inverse-Gamma}(a, 1/b) \\ \sigma^2 &\sim \text{Inverse-Gamma}(0.1, 0.1). \end{aligned}$$

I need the full conditionals

$$\begin{aligned} &\{\beta_1, \dots, \beta_p | \mathbf{Y}, \sigma^2, \tau_1^2, \dots, \tau_p^2, \lambda\}, \\ &\{\sigma^2 | \mathbf{Y}, \beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \lambda\}, \\ &\{\tau_1^2, \dots, \tau_p^2 | \mathbf{Y}, \beta_1, \dots, \beta_p, \sigma^2, \lambda\}, \\ &\{\lambda | \mathbf{Y}, \beta_1, \dots, \beta_p, \sigma^2, \tau_1^2, \dots, \tau_p^2\} \end{aligned}$$

which are all proportional to

$$p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) \times p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) \times p(\tau_1^2, \dots, \tau_p^2 | \lambda) p(\lambda) p(\sigma^2)$$

so I'll start with the posterior

$$\begin{aligned} p(\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda | \mathbf{Y}) &\propto p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) \\ &\quad \times p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) \\ &\quad \times p(\tau_1^2, \dots, \tau_p^2 | \lambda) p(\lambda) p(\sigma^2). \end{aligned}$$

As a function of just σ^2 , this is proportional to

$$\begin{aligned} &p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) p(\sigma^2) \\ &= N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) IG(a, b). \end{aligned}$$

Time to show this is inverse-gamma distributed.

$$\begin{aligned} &N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) IG(a, b) \\ &\propto (\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta) \right\} (\sigma^2)^{a-1} \exp \left\{ -\frac{b}{\sigma^2} \right\} \\ &= (\sigma^2)^{-(n/2+a)-1} \exp \left\{ -\frac{1}{\sigma^2} (2b + \frac{1}{2} (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)) \right\} \\ &= IG(n/2 + a, 2b + (\mathbf{y} - \mathbf{X}\beta)^\top (\mathbf{y} - \mathbf{X}\beta)/2) \end{aligned}$$

As a function of β , the conditional is proportional to

$$\begin{aligned} &p(\mathbf{Y}|\beta_1, \dots, \beta_p, \tau_1^2, \dots, \tau_p^2, \sigma^2, \lambda) p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) \\ &= N(\mathbf{X}\beta, \sigma^2 \mathbf{I}) \cdot \prod_{i=1}^p N(0, \tau_i^2) \\ &= N(\mathbf{X}\beta, \Sigma) \cdot N(0, \Omega), \text{ where } \Omega = \text{diag}(\tau_1^2, \dots, \tau_p^2) \\ &= N(m, \mathbf{M}) \end{aligned}$$

with $m = \mathbf{M}\mathbf{X}^\top \Sigma^{-1}y$ and $\mathbf{M} = (\mathbf{X}^\top \Sigma^{-1}\mathbf{X} + \Omega^{-1})^{-1}$. As a function of $\tau_1^2, \dots, \tau_p^2$, the conditional is non-standard, but it's proportional to

$$\begin{aligned} & p(\beta_1, \dots, \beta_p | \tau_1^2, \dots, \tau_p^2) p(\tau_1^2, \dots, \tau_p^2 | \lambda) \\ &= \prod_{i=1}^p N(0, \tau_i^2) \cdot \prod_{i=1}^p IG(1, \frac{\lambda^2}{2}) \end{aligned}$$

Finally, as a function of λ , it's proportional to

$$\begin{aligned} & p(\tau_1^2, \dots, \tau_p^2 | \lambda) p(\lambda) \\ &= \prod_{i=1}^p IG(1, \frac{\lambda^2}{2}) \cdot IG(0.1, 0.1) \end{aligned}$$

Now I can build an algorithm to iteratively update through these conditional distributions. I take the starting value of $\beta^{(0)}$ to be the least-squares solution $\hat{\beta}$.

Result: Samples from joint posterior $p(\beta, \sigma^2 | y)$

for s in # samples **do**

```

     $\sigma^{2(s+1)} \sim IG(n/2 + a, 2b + (\mathbf{y} - \mathbf{X}\beta^{(s)})^\top (\mathbf{y} - \mathbf{X}\beta^{(s)})/2)$ 
     $\beta^s \sim N_p(m, \mathbf{M})$ 
     $m^\top = \mathbf{M} = \tau_1^2, \dots, \tau_p^2 \sim N(\boldsymbol{\tau}^{(s)}, \delta \mathbf{I})^2$ 
     $\lambda^* \sim N(\lambda.s, \delta)^2$ 
     $\log(r) = \log p(\mathbf{y} | \beta^*, \tau_1^2, \dots, \tau_p^2, \sigma^{2(s)}) + \log p(\beta^* | \tau_1^2, \dots, \tau_p^2)$ 
                $- \log p(\mathbf{y} | \beta^{(s)}, \tau_1^2, \dots, \tau_p^2, \sigma^{2(s)}) - \log p(\beta^{(s)} | \tau_1^2, \dots, \tau_p^2)$ 
     $u \sim Unif(0, 1)$ 
    if  $\log(u) < \log(r)$  then
        |  $\boldsymbol{\tau}^{(s+1)} = \beta^*$ 
    else
        |  $\boldsymbol{\tau}^{(s+1)} = \beta^{(s)}$ 
    end

```

end

Algorithm 1: Gibbs and Metropolis

Part e.

temp

```
set.seed(1)
data("diabetes")
X <- cbind(diabetes$x); y <- diabetes$y; n <- nrow(X);
X <- cbind(rep(1,n),X); p <- ncol(X)
samples <- 5000; a <- b <- 0.1; delta <- 0.1
beta.s <- solve(t(X)%*%X)%*%t(X)%*%y
lambda.s <- 0.5
tau2.s <- rep(1,p)
s <- 1
beta<-sigma2<-rep(0,samples)
for(s in 1:samples){
  lambda.star <- rnorm(1,lambda2.s, 0.01)^2
  logr <-
    sum(dgamma(p,1,lambda.star/2, log = T))+
    sum(dgamma(1,0.1,10, log=T))-sum(dgamma(p,1,lambda.s/2, log=T))-
    sum(dgamma(1,0.1,10,log=T))
  if(log(runif(1))<logr){
    lambda.s <- lambda.star
  }else{
    lambda.s <- lambda.s
  }

  tau2.star <- rnorm(p,tau2.s, 0.01)^2
  logr <-
    dmvnorm(rep(0,p),mean=matrix(1,p,1),sigma=diag(p),log = T)+
    sum(dgamma(p,1,lambda2.s/2, log=T))-
    dmvnorm(rep(0,p),mean=rep(0,p),sigma=diag(tau2.s),log = T)-
    sum(dgamma(p,1,lambda2.s/2, log=T))
  if(log(runif(1))<logr){
    tau2.s <- tau2.star
  }else{
    tau2.s <- tau2.s
  }
  lambda <- 1/rgamma(1,0.1,10)
  tau2 <- 1/rgamma(p,1,lambda^2/2)
  sigma2.s <- rgamma(1,n/2+a, 1/(2*b+t(y-X)%*%beta.s)%*%(y-X)%*%beta.s)))
  M <- matrix(solve(1/sigma2.s*t(X)%*%X+diag(1/tau2)),p,p)
  m <- M%*%t(X)%*%y/sigma2.s
  beta.s <- t(rmvnorm(n=1,mean=m,sigma=diag(p)))
  sigma2 <- c(sigma2, sigma2.s)
  beta <- rbind(beta,t(matrix(beta.s)))
}

beta.ls <- solve(t(X)%*%X)%*%t(X)%*%y
fit <- glmnet(X[,-1], y)
df <- cbind(beta.ls,
colMeans(beta[floor(samples/4):samples,]), coef(fit,s=0))
colnames(df) <- c("Least Square", "Bayes", "glmnet")
df
```

Part f.

Implement such algorithm in R and compare your results with estimates obtained using **glmnet()**. In particular, you should test your results on the diabetes data available from lars, (use the matrix of predictors \mathbf{x}).

Part g.

Free λ and carry out a sensitivity analysis assessing the behavior of the posterior distribution of β and σ^2 , as hyper parameters a and b are changed. Explain clearly the rationale you use to assess sensitivity and provide recommendations for the analysis of the diabetes data.

Part h.

Implementation and benchmarking in Julia