

Homework 2

Ian Frankenburg

5/8/2020

Generalized Linear Models

Part (1).

Describe and implement a Metropolis-Hastings algorithm designed to obtain a MC with stationary distribution $p(\beta|Y)$

This MCMC implementation is pretty straightforward. I know my target $p(\beta|y)$ is proportional to

$$\left[\prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i} \right] \exp \left(-\frac{1}{2n} \beta^\top (X^\top X) \beta \right), \text{ where } p_i = \Phi(x_i^\top \beta).$$

My Metropolis algorithm will make symmetric proposals based around the current iteration of β with variance defined to be the inverse sample covariance $n(X^\top X)^{-1}$:

$$\beta^* \sim N(\beta^{(s)}, n(X^\top X)^{-1}).$$

```
require(survival)

## Loading required package: survival

data <- infert
target <- function(y, x, xtx, beta, n){
  p <- pnorm(x%*%beta)
  return(sum(y*log(p)+(1-y)*log(1-p))-1/(2*n)*t(beta)%*%xtx%*%beta)
}
samples <- 5000
y <- data$case
n <- length(y)
data$education <- as.numeric(as.factor(data$education))
edu1 <- data$education==2; edu2 <- data$education==3
x <- cbind(rep(1,n),data$age, data$parity, edu1, edu2, data$spontaneous, data$induced)
ncolx <- ncol(x)
xtx <- t(x)%*%x
xtx_inv <- solve(xtx)
beta <- t(rmvnorm(1,rep(0,ncol(x)),xtx_inv*n))
beta.chain <- matrix(0,nrow=ncolx,ncol=samples)
beta.chain[,1] <- beta
for(s in 2:samples){
  beta_new <- t(rmvnorm(1,as.vector(beta), xtx_inv))
  logr <- target(y,x,xtx,beta_new,n)-target(y,x,xtx,beta,n)
  if(log(runif(1)) < logr){beta <- beta_new}
```

```

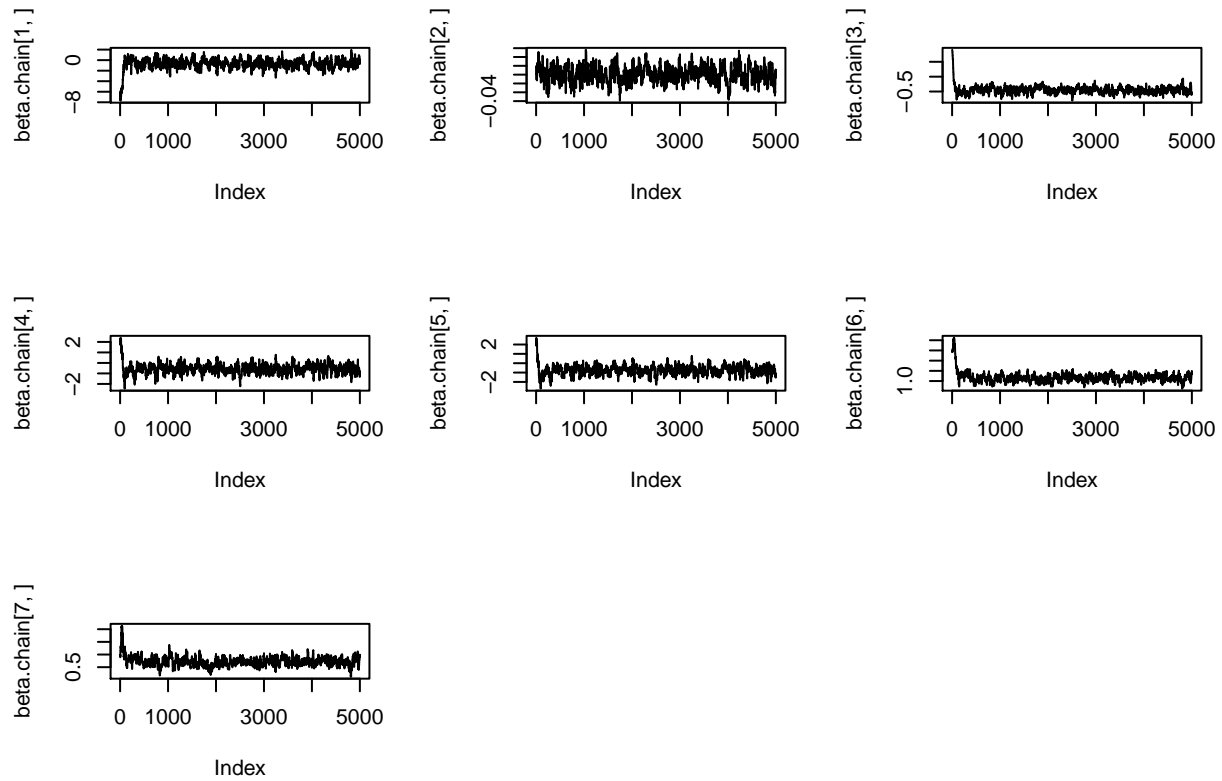
beta.chain[,s] <- beta
}

```

```

##      (Intercept)      x[, -1]      x[, -1] x[, -1]edu1 x[, -1]edu2      x[, -1]
## Bayes  -0.7085820  0.02092778 -0.4457638  -0.5614259  -0.7648931  1.176969
## GLM    -0.6413262  0.02056090 -0.4544283  -0.5705585  -0.7987104  1.173780
##      x[, -1]
## Bayes  0.7270190
## GLM    0.7214959

```



Part (2).

Describe and implement a data augmented (DA-MCMC) strategy targeting $p(\beta|Y)$

The original model is

$$P(y_i = 1|x_i, \beta) = \Phi(x_i^\top \beta).$$

This is equivalent to the model $P(y_i = 1|x_i, \beta) = P(\mathbf{1}(z_i > 0) = 1) = P(z_i > 0)$, where $z_i \stackrel{\text{iid}}{\sim} N(x_i^\top \beta, 1)$. This follows immediately since

$$\int_{-\infty}^{x_i^\top \beta} N(t; 0, 1) dt = \int_{-\infty}^0 N(z_i; x_i^\top \beta, 1) dz_i$$

by a change of variables $z_i := t - x_i^\top \beta$. Thus $\Phi(x_i^\top \beta) = P(z_i > 0)$.

In defining the latent model, the full conditionals $p(\beta|\mathbf{y}, \mathbf{z})$, $p(\mathbf{z}|\mathbf{y}, \beta)$ become tractible, so I can use a Gibbs sampler.

I'll start with $p(\beta|\mathbf{y}, \mathbf{z})$

$$\begin{aligned} p(\beta|\mathbf{y}, \mathbf{z}) &\propto p(\mathbf{y}, \mathbf{z}|\beta)p(\beta) \\ &= \overbrace{p(\mathbf{y}|\beta, \mathbf{z})}^{\cancel{p(\mathbf{y}|\beta, \mathbf{z})}} p(\mathbf{z}|\beta)p(\beta) = p(\mathbf{z}|\beta)p(\beta) \\ &= N(\mathbf{z}; X\beta, \mathbf{I})N(\beta; \mathbf{0}, n(X^\top X)^{-1}) \\ &= N\left(\frac{n}{n+1}(X^\top X)^{-1}X^\top \mathbf{y}, \frac{n}{n+1}(X^\top X)^{-1}\right) \end{aligned}$$

Now for $p(\mathbf{z}|\mathbf{y}, \beta)$.

$$\begin{aligned} p(\mathbf{z}|\mathbf{y}, \beta) &\propto p(\mathbf{y}, \beta, \mathbf{z}) \\ &\propto \overbrace{p(\mathbf{y}|\beta, \mathbf{z})}^{\cancel{p(\mathbf{y}|\beta, \mathbf{z})}} p(\mathbf{z}|\beta)\overbrace{p(\beta)}^{\cancel{p(\beta)}} \\ &\propto \prod_{i=1}^n [\mathbf{1}(y_i = 1)\mathbf{1}(z_i > 0) + \mathbf{1}(y_i = 0)\mathbf{1}(z_i < 0)]N(\mathbf{z}; X\beta, \mathbf{I}) \end{aligned}$$

Since our sampling model assumes the y 's are independent, so are the z 's and I can sample the full conditionals independently, i.e.

$$\begin{aligned} p(z_i|\mathbf{y}, \beta) &\propto p(\mathbf{y}, \beta, z_i) \propto p(\mathbf{y}|\beta, z_i)p(z_i|\beta) \\ &\propto [\mathbf{1}(y_i = 1)\mathbf{1}(z_i > 0) + \mathbf{1}(y_i = 0)\mathbf{1}(z_i < 0)]N(z_i; x_i^\top \beta, 1) \\ &= \begin{cases} N(z_i; x_i^\top \beta, 1) * \mathbf{1}_{[0, \infty)}(z_i) & \text{if } y_i = 1 \\ N(z_i; x_i^\top \beta, 1) * \mathbf{1}_{(-\infty, 0)}(z_i) & \text{if } y_i = 0 \end{cases} \end{aligned}$$

Now I can implement a Gibbs sampler to iteratively draw from these conditionals.

```
require(truncnorm)
```

```
## Loading required package: truncnorm
```

```
samples <- 5000
data <- infert
y <- data$case
n <- length(y)
data$education <- as.numeric(as.factor(data$education))
edu1 <- data$education==2; edu2 <- data$education==3
x <- cbind(rep(1,n), data$age, data$parity, edu1, edu2, data$spontaneous, data$induced)
p <- ncol(x)
txt <- t(x)%*%x
```

```

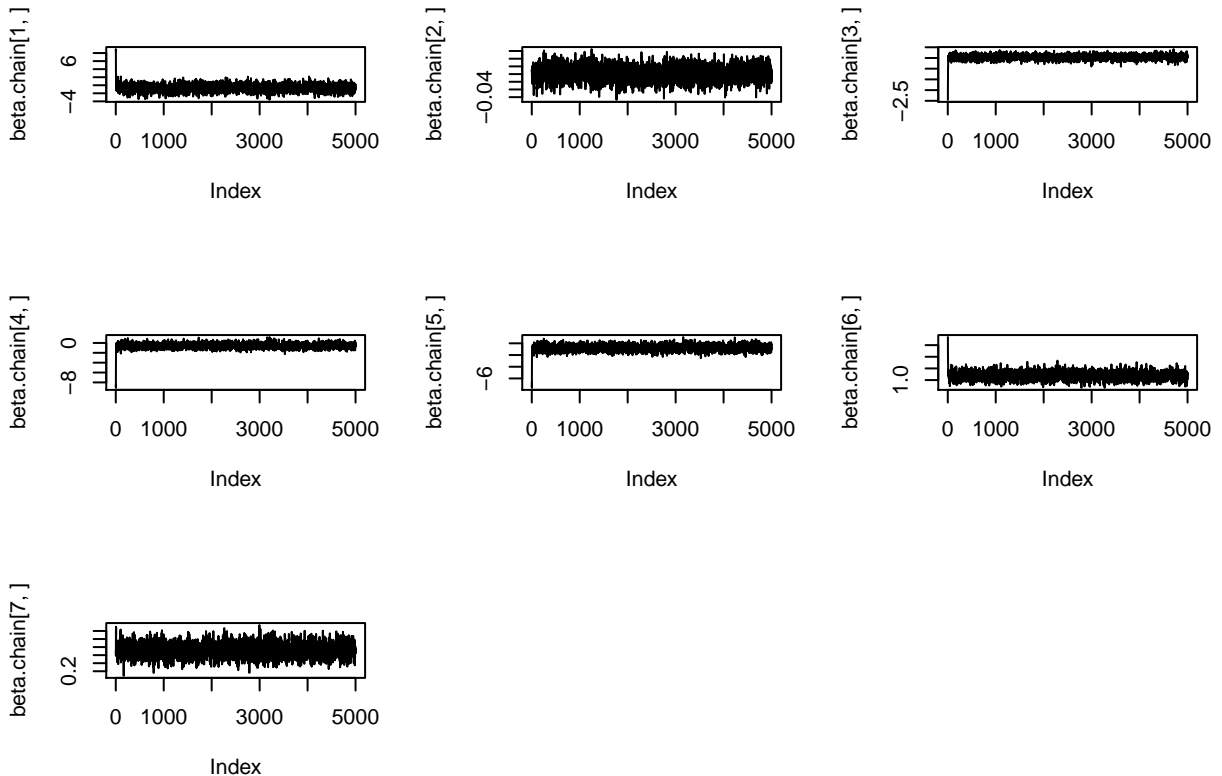
xtx_inv <- solve(xtx)
beta <- t(rmvnorm(1,rep(0,ncol(x)),xtx_inv*n))
beta.chain <- matrix(0,nrow=p,ncol=samples)
beta.chain[,1] <- beta
z <- t(rmvnorm(1,x%*%beta,sigma=diag(n)))
for(s in 2:samples){
  beta <- t(rmvnorm(1, n/(n+1)*xtx_inv%*%t(x)%*%z, n/(n+1)*xtx_inv))
  # truncated normal
  for(i in 1:n){
    if(y[i]==1){
      z[i] <- rtruncnorm(1, a = 0, b = Inf, mean = matrix(x[i,],ncol=p)%*%beta, sd=1)
    }else{
      z[i] <- rtruncnorm(1, a = -Inf, b = 0, mean = matrix(x[i,],ncol=p)%*%beta, sd=1)
    }
  }
  beta.chain[,s] <- beta
}

```

```

##      (Intercept)      x[, -1]      x[, -1] x[, -1]edu1 x[, -1]edu2 x[, -1]
## Bayes -0.6664524 0.02116593 -0.4601014 -0.5598672 -0.7926758 1.18323
## GLM   -0.6413262 0.02056090 -0.4544283 -0.5705585 -0.7987104 1.17378
##      x[, -1]
## Bayes 0.7268029
## GLM   0.7214959

```



Part (3).

Describe and implement a parameter expanded-data augmentation (PX-DA MCMC) algorithm targeting $p(\beta|Y)$.

For this parameter-expanded model, I'll introduce another parameter $\alpha^2 \sim IG(a, b)$ and I'll consider the transformation $w_i := \alpha z_i$. Then I'll use a Gibbs sampler to sample iteratively from the conditionals

$$\begin{aligned}
 p(\beta|\mathbf{y}, \mathbf{w}, \alpha) &\propto \prod_{i=1}^n N(w_i; \alpha x_i^\top \beta, \alpha^2) N(\beta; 0, n(X^\top X)^{-1}) \\
 &= N(Mm, M), \text{ where } M = \frac{n}{n+1}(X^\top X)^{-1} \text{ and } m = MX^\top \mathbf{w}/\alpha \\
 p(w_i|\mathbf{y}, \beta, \alpha) &= \begin{cases} N(w_i; \alpha x_i^\top \beta, \alpha^2) * \mathbf{1}_{[0, \infty)}(w_i) & \text{if } y_i = 1 \\ N(w_i; \alpha x_i^\top \beta, \alpha^2) * \mathbf{1}_{(-\infty, 0)}(w_i) & \text{if } y_i = 0 \end{cases} \\
 p(\alpha|\mathbf{y}, \beta, \mathbf{z}) &= \prod_{i=1}^n N(w_i; \alpha x_i^\top \beta, \alpha^2) IG(\alpha^2; a, b)
 \end{aligned}$$

```

require(truncnorm)
samples <- 5000
data <- infert
y <- data$case
data$education <- as.numeric(as.factor(data$education))
edu1 <- data$education==2; edu2 <- data$education==3
x <- cbind(rep(1,n),data$age, data$parity, edu1, edu2, data$spontaneous, data$induced)
n <- nrow(x); p <- ncol(x)
xtx <- t(x)%*%x
xtx_inv <- solve(xtx)
beta <- t(rmvnorm(1,rep(0,ncol(x)),xtx_inv*n))
beta.chain <- matrix(0,nrow=p,ncol=samples)
beta.chain[,1] <- beta
alpha2 <- 1
z <- t(rmvnorm(1,x)%*%beta,sigma=alpha2 * diag(n)))
for(s in 2:samples){
  beta <- t(rmvnorm(1, n/(n+1)*xtx_inv%*%t(x)%*%z%*%solve(sqrt(alpha2)), n/(n+1)*xtx_inv))
  # truncated normal
  for(i in 1:n){
    if(y[i]==1){
      z[i] <- rtruncnorm(1, a = 0, b = Inf, mean = matrix(x[i,],ncol=p)%*%beta, sd=1)
    }else{
      z[i] <- rtruncnorm(1, a = -Inf, b = 0, mean = matrix(x[i,],ncol=p)%*%beta, sd=1)
    }
  }
  rss <- t((z-x)%*%beta)%*%(z-x)%*%beta)/2
  d <- rchisq(1, df=n)
  alpha2 <- sqrt(rss/d)
  beta.chain[,s] <- beta
}








```

```

##      (Intercept)      x[, -1]      x[, -1] x[, -1]edu1 x[, -1]edu2      x[, -1]
## Bayes -0.8661613 0.02697548 -0.7355791 -0.8758183 -1.2404586 1.847926
## GLM   -0.6413262 0.02056090 -0.4544283 -0.5705585 -0.7987104 1.173780
##      x[, -1]

```

```
## Bayes 1.2268688
## GLM 0.7214959
```



Part (6).

For logit model, describe and a random walk MH targeting $p(\beta|Y)$

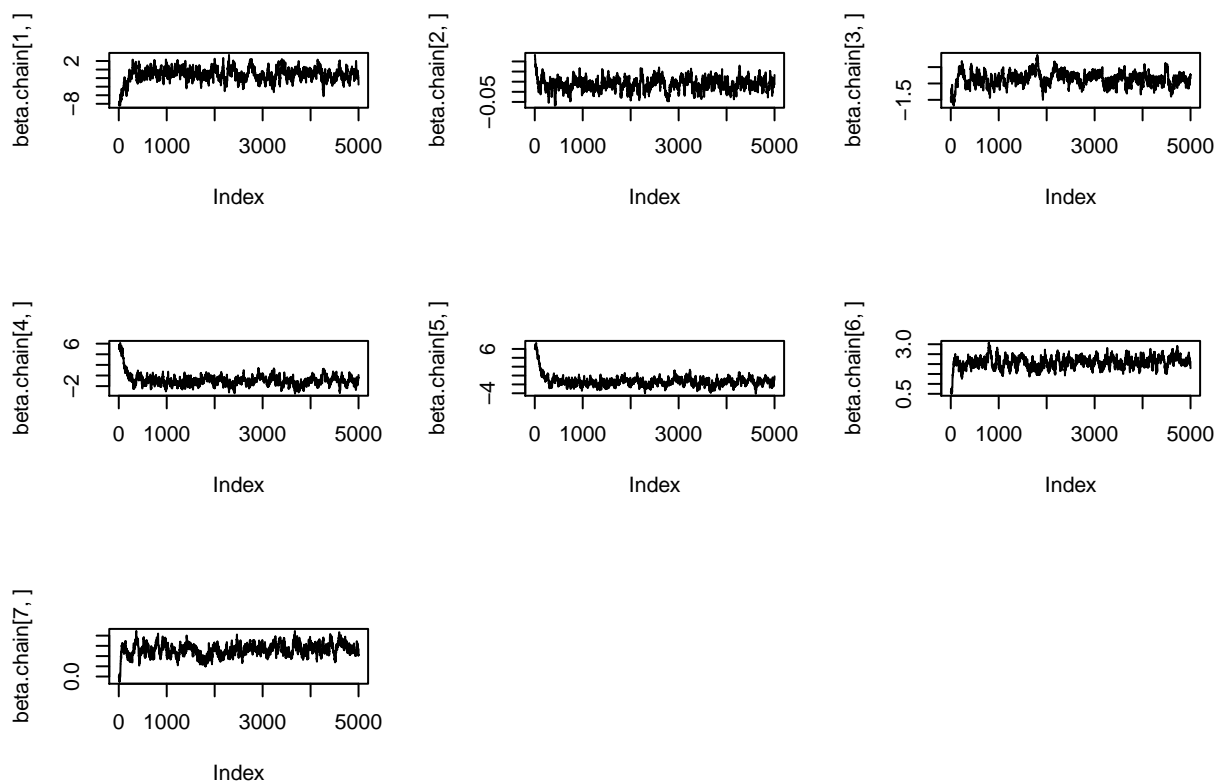
Similar to the probit case, the posterior is intractible but can be written as

$$\begin{aligned} p(\beta|y) &\propto \prod_{i=1}^n p_i^{y_i} (1-p_i)^{1-y_i} p(\beta) \\ &= \prod_{i=1}^n \left[\frac{e^{x_i^\top \beta}}{1 + e^{x_i^\top \beta}} \right]^{y_i} \left[1 - \frac{e^{x_i^\top \beta}}{1 + e^{x_i^\top \beta}} \right]^{1-y_i} N(\beta; 0, n(X^\top X)^{-1}) \\ &= \prod_{i=1}^n \frac{(e^{x_i^\top \beta})^{y_i}}{1 + e^{x_i^\top \beta}} N(\beta; 0, n(X^\top X)^{-1}) \end{aligned}$$

This will be my target in the random walk MH.

```
data <- infert
y <- data$case
data$education <- as.numeric(as.factor(data$education))
edu1 <- data$education==2; edu2 <- data$education==3
x <- cbind(rep(1,n),data$age, data$parity, edu1, edu2, data$spontaneous, data$induced)
target <- function(y, x, xtx, beta, n){
  p <- exp(x%*%beta)/(1+exp(x%*%beta))
  return(sum(y*log(p)+(1-y)*log(1-p))-1/(2*n)*t(beta)%*%xtx%*%beta)
}
samples <- 5000
n <- nrow(x); ncolx <- ncol(x)
xtx <- t(x)%*%x
xtx_inv <- solve(xtx)
beta <- t(rmvnorm(1,rep(0,ncol(x)),xtx_inv*n))
beta.chain <- matrix(0,nrow=ncolx,ncol=samples)
beta.chain[,1] <- beta
for(s in 2:samples){
  beta_new <- t(rmvnorm(1,as.vector(beta), xtx_inv))
  logr <- target(y,x,xtx,beta_new,n)-target(y,x,xtx,beta,n)
  if(log(runif(1)) < logr){beta <- beta_new}
  beta.chain[,s] <- beta
}
```

```
##      (Intercept)      x[, -1]      x[, -1] x[, -1]edu1 x[, -1]edu2      x[, -1]
## Bayes   -1.175232  0.03751771 -0.8646845  -0.8969633   -1.304649  2.077206
## GLM     -1.149237  0.03958200 -0.8282774  -1.0442436   -1.403205  2.045905
##      x[, -1]
## Bayes  1.303842
## GLM    1.288757
```



Part (7).

Describe and implement a Langevin-Hastings algorithm designed to obtain a MC with stationary distribution $p(\beta|Y)$

The Langevin-Hastings algorithm will work by utilizing the a 2nd order Taylor approximation of the target distribution. That way I can try to match the proposal to the target.

If I was proposing for β at iteration s , I'd use something like

$$\beta^* \sim N\left(\beta^{(s)} + \frac{1}{2}\sigma^2 \nabla \log(p(\mathbf{y}|\beta^{(s)})p(\beta^{(s)})), \sigma^2 \mathbf{I}\right)$$

Earlier I showed the poserior is proportional to

$$\prod_{i=1}^n \frac{(e^{x_i^\top \beta})^{y_i}}{1 + e^{x_i^\top \beta}} N(\beta; 0, n(X^\top X)^{-1})$$

so the log-posterior is







$$\begin{aligned} & \sum_{i=1}^n y_i x_i^\top \beta - \log(1 + e^{x_i^\top \beta}) - \frac{1}{2n} \beta^\top (X^\top X) \beta \\ \Rightarrow & \nabla \left\{ \sum_{i=1}^n y_i x_i^\top \beta - \log(1 + e^{x_i^\top \beta}) - \frac{1}{2n} \beta^\top (X^\top X) \beta \right\} \\ = & \sum_{i=1}^n \left[y_i - \frac{e^{x_i^\top \beta}}{1 + e^{x_i^\top \beta}} \right] x_i^\top - \frac{1}{n} \beta^\top (X^\top X) \end{aligned}$$

```
gradient <- function(y, x, xtx, beta, p){
  n <- length(y)
  return(t(y-p)%*%x-1/n*t(beta)%*%xtx)
}
target <- function(y, x, xtx, beta, n){
  return(sum(y*log(p)+(1-y)*log(1-p))-1/(2*n)*t(beta)%*%xtx%*%beta)
}
data <- infert
y <- data$case
data$education <- as.numeric(as.factor(data$education))
edu1 <- data$education==2; edu2 <- data$education==3
x <- cbind(rep(1,n),data$age, data$parity, edu1, edu2, data$spontaneous, data$induced)
samples <- 5000
n <- nrow(x); ncolx <- ncol(x)
xtx <- t(x)%*%x
xtx_inv <- solve(xtx)
beta <- t(rmvnorm(1,rep(0,ncolx),xtx_inv*n))
beta.chain <- matrix(0,nrow=ncolx,ncol=samples)
beta.chain[,1] <- beta
c <- 0.005
num_accepted <- accept_ratio <- 0
for(s in 2:samples){
  p <- exp(x%*%beta)/(1+exp(x%*%beta))
  grad <- gradient(y, x, xtx, beta, p)
  beta_new <- t(rmvnorm(1,as.vector(beta) + c^2*as.vector(grad), c*diag(ncolx)))
  logr <- target(y,x,xtx,beta_new,n)-target(y,x,xtx,beta,n)
  if(log(runif(1)) < logr){
```

```

    beta <- beta_new
    num_accepted <- num_accepted+1
  }
  accept_ratio <- num_accepted/s
  beta.chain[,s] <- beta
}


```



```

##      (Intercept)      x[, -1]      x[, -1] x[, -1]edu1 x[, -1]edu2      x[, -1]
## Bayes  0.3153487 -0.07770807  0.6285502   1.4391816   2.5141209  0.2328441
## GLM   -0.6413262  0.02056090 -0.4544283  -0.5705585  -0.7987104  1.1737802
##      x[, -1]
## Bayes -2.3075979
## GLM   0.7214959

```



These chains look crazy, but I'm not sure what's happening. I think the tuning parameter c is the culprit.

Part (7).

Describe and implement an adaptive Metropolis-Hastings algorithm designed to obtain a MC with stationary distribution $p(\beta|Y)$

I think in this in this problem, I'll make proposals at step s such as

$$\beta^* \sim \beta^{(s)} + N(0, c\Sigma^{(s)})$$

and adaptively change c based on my acceptance ratio.

```
data <- infert
y <- data$case
data$education <- as.numeric(as.factor(data$education))
edu1 <- data$education==2; edu2 <- data$education==3
x <- cbind(rep(1,n),data$age, data$parity, edu1, edu2, data$spontaneous, data$induced)
target <- function(y, x, xtx, beta, n){
  p <- exp(x%*%beta)/(1+exp(x%*%beta))
  return(sum(y*log(p)+(1-y)*log(1-p))-1/(2*n)*t(beta)%*%xtx%*%beta)
}
samples <- 5000
n <- nrow(x); ncolx <- ncol(x)
xtx <- t(x)%*%x
xtx_inv <- solve(xtx)
beta <- t(rmvnorm(1,rep(0,ncol(x)),xtx_inv*n))
beta.chain <- matrix(0,nrow=ncolx,ncol=samples)
beta.chain[,1] <- beta
num_accepted <- accept_ratio <- 0
c <- 10
for(s in 2:samples){
  beta_new <- t(rmvnorm(1,as.vector(beta), c*xtx_inv))
  logr <- target(y,x,xtx,beta_new,n)-target(y,x,xtx,beta,n)
  if(log(runif(1)) < logr){
    beta <- beta_new
    num_accepted <- num_accepted+1
  }
  accept_ratio <- num_accepted/s
  if(accept_ratio > 0.4){
    c <- c+1
  }else{
    c <- c/2
  }
  beta.chain[,s] <- beta
}
```

```
##      (Intercept)      x[, -1]      x[, -1] x[, -1]edu1 x[, -1]edu2      x[, -1]
## Bayes   -1.137631  0.03744315 -0.7925801  -0.9782338   -1.341475  2.014576
## GLM     -1.149237  0.03958200 -0.8282774  -1.0442436   -1.403205  2.045905
##      x[, -1]
## Bayes  1.203220
## GLM    1.288757
```

